

## Лінійні моделі

```
In [1]: import os
import sklearn
import numpy as np
import matplotlib.pyplot as plt

%matplotlib inline
```


### Лінійна регресія

Згенеруємо синтетичні дані для прикладу

```
In [2]: np.random.seed(42)
```

```
In [3]: X = 2 * np.random.rand(100, 1)
y = 4 + 3 * X + np.random.randn(100, 1)
```

```
In [4]: plt.plot(X, y, "b.")
plt.xlabel("$x_1$", fontsize=18)
plt.ylabel("$y$", rotation=0, fontsize=18)
plt.axis([0, 2, 0, 15])
plt.show()
```



```
In [5]: X_b = np.c_[np.ones((100, 1)), X] # add x0 = 1 to each instance
theta_best = np.linalg.inv(X_b.T.dot(X_b)).dot(X_b.T).dot(y)
```

Знайдені коефіцієнти

```
In [6]: theta_best
```

```
Out[6]: array([[4.21509616],
               [2.77011339]])
```


Як отримати значення функції регресії у новій точці?

```
In [7]: X_new = np.array([[0], [2]])
X_new_b = np.c_[np.ones((2, 1)), X_new] # add x0 = 1 to each instance
y_predict = X_new_b.dot(theta_best)
y_predict
```

```
Out[7]: array([[4.21509616],
               [9.75532293]])
```

Судячи з графіку, виглядає логічно.

```
In [8]: plt.plot(X_new, y_predict, "r-")
plt.plot(X, y, "b.")
plt.axis([0, 2, 0, 15])
plt.show()
```



Розв'яжемо цю ж задачу за допомогою sklearn

```
In [9]: from sklearn.linear_model import LinearRegression

lin_reg = LinearRegression()
lin_reg.fit(X, y)
lin_reg.intercept_, lin_reg.coef_
```

```
Out[9]: (array([4.21509616]), array([[2.77011339]]))
```

## Поліноміальна регресія


```
In [10]: import numpy as np
import numpy.random as rnd

np.random.seed(42)
```

Згенеруємо штучний датасет з відгуком y, що поліноміально залежить від ознак X.

```
In [11]: m = 100
X = 6 * np.random.rand(m, 1) - 3
y = 0.5 * X**2 + X + 2 + np.random.randn(m, 1)
```

```
In [12]: plt.plot(X, y, "b.")
plt.xlabel("$x_1$", fontsize=18)
plt.ylabel("$y$", rotation=0, fontsize=18)
plt.axis([-3, 3, 0, 10])
plt.show()
```



"Доповнимо" датасет поліноміальними ознаками

```
In [13]: from sklearn.preprocessing import PolynomialFeatures
poly_features = PolynomialFeatures(degree=2, include_bias=False)
X_poly = poly_features.fit_transform(X)
print(X[0])
print(X_poly[0])
print(X[0]**2)
```

```
[-0.75275929]
[-0.75275929  0.56664654]
[0.56664654]
```

```
In [14]: lin_reg = LinearRegression()
lin_reg.fit(X_poly, y)
lin_reg.intercept_, lin_reg.coef_
```

```
Out[14]: (array([1.78134581]), array([[0.93366893, 0.56456263]]))
```

Давайте спробуємо різні степені поліноміальних ознак: 1, 2, 300

```


In [15]: from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import Pipeline

X_new = np.linspace(-3, 3, num=200).reshape(-1,1)

for style, width, degree in (("g-", 1, 300), ("b--", 2, 2), ("r+", 2, 1)):
    polybig_features = PolynomialFeatures(degree=degree, include_bias=False)
    std_scaler = StandardScaler()
    lin_reg = LinearRegression()
    polynomial_regression = Pipeline([
        ("poly_features", polybig_features),
        ("std_scaler", std_scaler),
        ("lin_reg", lin_reg),
    ])
    polynomial_regression.fit(X, y)
    y_newbig = polynomial_regression.predict(X_new)
    plt.plot(X_new, y_newbig, style, label=str(degree), linewidth=width)

plt.plot(X, y, "b.", linewidth=3)
plt.legend(loc="upper left")
plt.xlabel("$x_1$", fontsize=18)
plt.ylabel("$y$", rotation=0, fontsize=18)
plt.axis([-3, 3, 0, 10])
plt.show()

```



### Середньоквадратична помилка моделі

```

In [16]: from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error

```

```

In [17]: np.random.seed(42)
m = 100
X = 6 * np.random.rand(m, 1) - 3
y = 2 + X + 0.5 * X**2 + np.random.randn(m, 1)

```

```

In [18]: X_train, X_val, y_train, y_val = train_test_split(X, y.ravel(), test_size=0.5, random_state=10)

```

```

In [19]: X_train.shape

```

```

Out[19]: (50, 1)

```

```
In [20]: poly_scaler = Pipeline([
          ("poly_features", PolynomialFeatures(degree=5, include_bias=False)),
          ("std_scaler", StandardScaler())
        ])

X_train_poly_scaled = poly_scaler.fit_transform(X_train)
X_val_poly_scaled = poly_scaler.transform(X_val)
```

```
In [21]: model = LinearRegression()
         model.fit(X_train_poly_scaled, y_train)
```

```
Out[21]: LinearRegression()
```

```
In [22]: y_val_predict = model.predict(X_val_poly_scaled)
         val_error = mean_squared_error(y_val, y_val_predict)
```

```
In [23]: print(val_error)

0.9418567491761436
```

```
In [24]: def get_error(poly_degree, X_train, X_val, y_train, y_val):
         poly_scaler = Pipeline([
           ("poly_features", PolynomialFeatures(degree=poly_degree, include_bias=False)),
           ("std_scaler", StandardScaler())
         ])

         X_train_poly_scaled = poly_scaler.fit_transform(X_train)
         X_val_poly_scaled = poly_scaler.transform(X_val)

         model = LinearRegression()
         model.fit(X_train_poly_scaled, y_train)
         y_val_predict = model.predict(X_val_poly_scaled)


         val_error = mean_squared_error(y_val, y_val_predict)

         return val_error
```

```
In [25]: val_errors = [get_error(poly_degree, X_train, X_val, y_train, y_val) for poly_degree in range(1,11)]
```

```
In [26]: plt.plot(list(range(1,11)), np.sqrt(val_errors), "b-", linewidth=3, label="val")
```

```
Out[26]: [<matplotlib.lines.Line2D at 0x7f88eeeaab50>]
```




```
In [27]: final_pipeline = poly_scaler = Pipeline([
    ("poly_features", PolynomialFeatures(degree=2, include_bias=False)),
    ("std_scaler", StandardScaler()),
    ("linear_model", LinearRegression())
])
```

```
In [28]: final_pipeline.fit(X,y)
```

```
Out[28]: Pipeline(steps=[('poly_features', PolynomialFeatures(include_bias=False)),
    ('std_scaler', StandardScaler()),
    ('linear_model', LinearRegression())])
```

```
In [29]: X_test = np.linspace(-3, 3, 100)
plt.plot(X_test, final_pipeline.predict(X_test[:, np.newaxis]), label="Model")
plt.scatter(X, y, label="True function")
```

```
Out[29]: <matplotlib.collections.PathCollection at 0x7f88eee8fbd0>
```



## Регуляризація

```
In [30]: np.random.seed(42)
m = 20
X = 3 * np.random.rand(m, 1)
y = 1 + 0.5 * X + np.random.randn(m, 1) / 1.5
X_new = np.linspace(0, 3, 100).reshape(100, 1)
```

```
In [31]: from sklearn.linear_model import Ridge
ridge_reg = Ridge(alpha=1, random_state=42)
ridge_reg.fit(X, y)
ridge_reg.predict([[1.5]])
```

```
Out[31]: array([[1.55071465]])
```

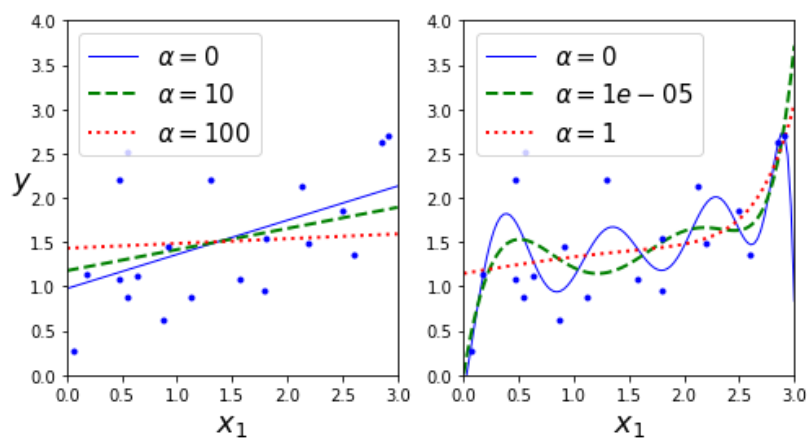
```

In [32]: def plot_model(model_class, polynomial, alphas, degree = 10, **model_kargs):
    for alpha, style in zip(alphas, ("b-", "g--", "r:")):
        model = model_class(alpha, **model_kargs) if alpha > 0 else LinearRe
gression()
        if polynomial:
            model = Pipeline([
                ("poly_features", PolynomialFeatures(degree=degree, incl
ude_bias=False)),
                ("std_scaler", StandardScaler()),
                ("regul_reg", model),
            ])
            model.fit(X, y)
            y_new_regul = model.predict(X_new)
            lw = 2 if alpha > 0 else 1
            plt.plot(X_new, y_new_regul, style, linewidth=lw, label=r"$\alpha =
{}$".format(alpha))
            plt.plot(X, y, "b.", linewidth=3)
            plt.legend(loc="upper left", fontsize=15)
            plt.xlabel("$x_1$", fontsize=18)
            plt.axis([0, 3, 0, 4])

plt.figure(figsize=(8,4))
plt.subplot(121)
plot_model(Ridge, polynomial=False, alphas=(0, 10, 100), random_state=42)
plt.ylabel("$y$", rotation=0, fontsize=18)
plt.subplot(122)
plot_model(Ridge, polynomial=True, alphas=(0, 10**-5, 1), random_state=42)

plt.show()


```



```
In [33]: from sklearn.linear_model import Lasso


plt.figure(figsize=(8,4))
plt.subplot(121)
plot_model(Lasso, polynomial=False, alphas=(0, 0.1, 1), random_state=42)
plt.ylabel("$y$", rotation=0, fontsize=18)
plt.subplot(122)
plot_model(Lasso, polynomial=True, alphas=(0, 10**-7, 1), random_state=42)
plt.show()
```

/home/daryna/anaconda3/envs/ml\_ukma/lib/python3.7/site-packages/sklearn/linear\_model/\_coordinate\_descent.py:531: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 2.802867703827423, tolerance: 0.0009294783355207351 positive)



## Логістична регресія

```
In [34]: t = np.linspace(-10, 10, 100)
sig = 1 / (1 + np.exp(-t))
plt.figure(figsize=(9, 3))
plt.plot([-10, 10], [0, 0], "k-")
plt.plot([-10, 10], [0.5, 0.5], "k:")
plt.plot([-10, 10], [1, 1], "k:")
plt.plot([0, 0], [-1.1, 1.1], "k-")
plt.plot(t, sig, "b-", linewidth=2, label=r"$\sigma(t) = \frac{1}{1 + e^{-t}}$")
plt.xlabel("t")
plt.legend(loc="upper left", fontsize=20)
plt.axis([-10, 10, -0.1, 1.1])
plt.show()
```





```
In [35]: from sklearn import datasets  
iris = datasets.load_iris()  
list(iris.keys())
```

```
Out[35]: ['data',  
          'target',  
          'frame',  
          'target_names',  
          'DESCR',  
          'feature_names',  
          'filename']
```

```
In [36]: print(iris.DESCR)
```

```
.. _iris_dataset:
```

```
Iris plants dataset
```

```
-----
```

```
**Data Set Characteristics:**
```

```
:Number of Instances: 150 (50 in each of three classes)
:Number of Attributes: 4 numeric, predictive attributes and the class
:Attribute Information:
  - sepal length in cm
  - sepal width in cm
  - petal length in cm
  - petal width in cm
  - class:
    - Iris-Setosa
    - Iris-Versicolour
    - Iris-Virginica
```

```
:Summary Statistics:
```

	Min	Max	Mean	SD	Class Correlation
sepal length:	4.3	7.9	5.84	0.83	0.7826
sepal width:	2.0	4.4	3.05	0.43	-0.4194
petal length:	1.0	6.9	3.76	1.76	0.9490 (high!)
petal width:	0.1	2.5	1.20	0.76	0.9565 (high!)

```
:Missing Attribute Values: None
:Class Distribution: 33.3% for each of 3 classes.
:Creator: R.A. Fisher
:Donor: Michael Marshall (MARSHALL%PLU@io.arc.nasa.gov)
:Date: July, 1988
```

The famous Iris database, first used by Sir R.A. Fisher. The dataset is taken from Fisher's paper. Note that it's the same as in R, but not as in the UCI Machine Learning Repository, which has two wrong data points.

This is perhaps the best known database to be found in the pattern recognition literature. Fisher's paper is a classic in the field and is referenced frequently to this day. (See Duda & Hart, for example.) The data set contains 3 classes of 50 instances each, where each class refers to a type of iris plant. One class is linearly separable from the other 2; the latter are NOT linearly separable from each other.

```
.. topic:: References
```

- Fisher, R.A. "The use of multiple measurements in taxonomic problems" Annual Eugenics, 7, Part II, 179-188 (1936); also in "Contributions to Mathematical Statistics" (John Wiley, NY, 1950).
- Duda, R.O., & Hart, P.E. (1973) Pattern Classification and Scene Analysis. (Q327.D83) John Wiley & Sons. ISBN 0-471-22361-1. See page 218.
- Dasarathy, B.V. (1980) "Nosing Around the Neighborhood: A New System Structure and Classification Rule for Recognition in Partially Exposed Environments". IEEE Transactions on Pattern Analysis and Machine Intelligence, Vol. PAMI-2, No. 1, 67-71.
- Gates, G.W. (1972) "The Reduced Nearest Neighbor Rule". IEEE Transactions on Information Theory, May 1972, 431-433.
- See also: 1988 MLC Proceedings, 54-64. Cheeseman et al's AUTOCLASS II conceptual clustering system finds 3 classes in the data.
- Many, many more ...

```
In [37]: iris["data"].shape
```

```
Out[37]: (150, 4)
```


```
In [38]: X = iris["data"][:, 3:] # petal width  
y = (iris["target"] == 2).astype(np.int) # 1 if Iris virginica, else 0
```

```
In [39]: from sklearn.linear_model import LogisticRegression  
log_reg = LogisticRegression(random_state=42)  
log_reg.fit(X, y)
```

```
Out[39]: LogisticRegression(random_state=42)
```

```
In [40]: X_new = np.linspace(0, 3, 1000).reshape(-1, 1)  
y_proba = log_reg.predict_proba(X_new)  
  
plt.plot(X_new, y_proba[:, 1], "g-", linewidth=2, label="Iris virginica")  
plt.plot(X_new, y_proba[:, 0], "b--", linewidth=2, label="Not Iris virginica")
```


```
Out[40]: [<matplotlib.lines.Line2D at 0x7f88ee1286d0>]
```



```
In [41]: X_new = np.linspace(0, 3, 1000).reshape(-1, 1)
y_proba = log_reg.predict_proba(X_new)
decision_boundary = X_new[y_proba[:, 1] >= 0.5][0]

plt.figure(figsize=(8, 3))
plt.plot(X[y==0], y[y==0], "bs")
plt.plot(X[y==1], y[y==1], "g^")
plt.plot([decision_boundary, decision_boundary], [-1, 2], "k:", linewidth=2)
plt.plot(X_new, y_proba[:, 1], "g-", linewidth=2, label="Iris virginica")
plt.plot(X_new, y_proba[:, 0], "b--", linewidth=2, label="Not Iris virginic
a")
plt.text(decision_boundary+0.02, 0.15, "Decision boundary", fontsize=14, co
lor="k", ha="center")
plt.arrow(decision_boundary, 0.08, -0.3, 0, head_width=0.05, head_length=0.
1, fc='b', ec='b')
plt.arrow(decision_boundary, 0.92, 0.3, 0, head_width=0.05, head_length=0.1,
fc='g', ec='g')
plt.xlabel("Petal width (cm)", fontsize=14)
plt.ylabel("Probability", fontsize=14)
plt.legend(loc="center left", fontsize=14)
plt.axis([0, 3, -0.02, 1.02])
plt.show()
```

/home/daryna/anaconda3/envs/ml\_ukma/lib/python3.7/site-packages/matplotlib/pa  
tches.py:1338: VisibleDeprecationWarning: Creating an ndarray from ragged nes  
ted sequences (which is a list-or-tuple of lists-or-tuples-or ndarrays with d  
ifferent lengths or shapes) is deprecated. If you meant to do this, you must  
specify 'dtype=object' when creating the ndarray  
verts = np.dot(coords, M) + (x + dx, y + dy)



```
In [42]: decision_boundary
```

```
Out[42]: array([1.66066066])
```

```
In [43]: from sklearn.linear_model import LogisticRegression

X = iris["data"][:, (2, 3)] # petal length, petal width
y = (iris["target"] == 2).astype(np.int)

log_reg = LogisticRegression(solver="lbfgs", C=10**10, random_state=42)
log_reg.fit(X, y)

x0, x1 = np.meshgrid(
    np.linspace(2.9, 7, 500).reshape(-1, 1),
    np.linspace(0.8, 2.7, 200).reshape(-1, 1),
)
X_new = np.c_[x0.ravel(), x1.ravel()]


y_proba = log_reg.predict_proba(X_new)

plt.figure(figsize=(10, 4))
plt.plot(X[y==0, 0], X[y==0, 1], "bs")
plt.plot(X[y==1, 0], X[y==1, 1], "g^")

zz = y_proba[:, 1].reshape(x0.shape)
contour = plt.contour(x0, x1, zz, cmap=plt.cm.brg)

left_right = np.array([2.9, 7])
boundary = -(log_reg.coef_[0][0] * left_right + log_reg.intercept_[0]) / log_reg.coef_[0][1]

plt.clabel(contour, inline=1, fontsize=12)
plt.plot(left_right, boundary, "k--", linewidth=3)
plt.text(3.5, 1.5, "Not Iris virginica", fontsize=14, color="b", ha="center")
plt.text(6.5, 2.3, "Iris virginica", fontsize=14, color="g", ha="center")
plt.xlabel("Petal length", fontsize=14)
plt.ylabel("Petal width", fontsize=14)
plt.axis([2.9, 7, 0.8, 2.7])
plt.show()
```



```
In [44]: X = iris["data"][:, (2, 3)] # petal length, petal width
y = iris["target"]

softmax_reg = LogisticRegression(multi_class="multinomial", solver="lbfgs", C=10, random_state=42)
softmax_reg.fit(X, y)
```

Out[44]: LogisticRegression(C=10, multi\_class='multinomial', random\_state=42)

```

In [45]: x0, x1 = np.meshgrid(
            np.linspace(0, 8, 500).reshape(-1, 1),
            np.linspace(0, 3.5, 200).reshape(-1, 1),
        )
X_new = np.c_[x0.ravel(), x1.ravel()]

y_proba = softmax_reg.predict_proba(X_new)
y_predict = softmax_reg.predict(X_new)


zz1 = y_proba[:, 1].reshape(x0.shape)
zz = y_predict.reshape(x0.shape)

plt.figure(figsize=(10, 4))
plt.plot(X[y==2, 0], X[y==2, 1], "g^", label="Iris virginica")
plt.plot(X[y==1, 0], X[y==1, 1], "bs", label="Iris versicolor")
plt.plot(X[y==0, 0], X[y==0, 1], "yo", label="Iris setosa")

from matplotlib.colors import ListedColormap
custom_cmap = ListedColormap(['#fafab0', '#9898ff', '#a0faa0'])

plt.contourf(x0, x1, zz, cmap=custom_cmap)
contour = plt.contour(x0, x1, zz1, cmap=plt.cm.brg)
plt.clabel(contour, inline=1, fontsize=12)
plt.xlabel("Petal length", fontsize=14)
plt.ylabel("Petal width", fontsize=14)
plt.legend(loc="center left", fontsize=14)
plt.axis([0, 7, 0, 3.5])
plt.show()

```



Нарешті, натренуємо модель на усіх ознаках.

```

In [46]: X = iris["data"]
y = iris["target"]

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.15)

```

```
In [47]: softmax_reg = LogisticRegression(multi_class="multinomial", solver="lbfgs", C
        =10, random_state=42)
        softmax_reg.fit(X_train, y_train)

/home/daryna/anaconda3/envs/ml_ukma/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:764: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
    https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
    https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
    extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)

Out[47]: LogisticRegression(C=10, multi_class='multinomial', random_state=42)

In [48]: y_pred = softmax_reg.predict(X_test)

In [49]: from sklearn.metrics import accuracy_score, confusion_matrix

In [50]: accuracy_score(y_test, y_pred)

Out[50]: 1.0

In [51]: confusion_matrix(y_test, y_pred)

Out[51]: array([[9, 0, 0],
               [0, 8, 0],
               [0, 0, 6]])
```

## Завдання

```
In [52]: import pandas as pd

        from IPython.display import Image

        from sklearn.linear_model import LinearRegression, Lasso, Ridge, LogisticRegression
        from sklearn.preprocessing import PolynomialFeatures, StandardScaler
        from sklearn.metrics import mean_squared_error
        from sklearn.model_selection import train_test_split
```

Завдання 1. Знайдіть у явному вигляді розв'язок задачі лінійної регресії з квадратичною регуляризацією (Ridge). Сфотографуйте виведення та підв'яжіть до цього файлу (``)



In [53]: `PICTURE_PATH = "../images/"`  
`Image(filename = PICTURE_PATH + "02-ridge.jpg", width=900, height=1600)`

Out[53]:

Ridge (розб'язок)

$$L(X, \vec{\beta}) = \|y - X\vec{\beta}\|_2^2 + \lambda \|\vec{\beta}\|_2^2 =$$

$$= (\vec{\beta}^T X^T y - y^T X \vec{\beta} - \vec{\beta}^T X^T y + \vec{\beta}^T X^T X \vec{\beta} + \lambda \vec{\beta}^T \vec{\beta}) =$$

$$= y^T y - y^T X \vec{\beta} - \vec{\beta}^T X^T y + \vec{\beta}^T X^T X \vec{\beta} + \lambda \vec{\beta}^T \vec{\beta}$$

$$\frac{\partial L(X, \vec{\beta})}{\partial \vec{\beta}} = \frac{\partial (y^T y - y^T X \vec{\beta} - \vec{\beta}^T X^T y + \vec{\beta}^T X^T X \vec{\beta} + \lambda \vec{\beta}^T \vec{\beta})}{\partial \vec{\beta}} =$$

$$= -2 y^T X + 2 \vec{\beta}^T X^T X + 2 \lambda \vec{\beta}$$

$$-2 y^T X + 2 \vec{\beta}^T X^T X + 2 \lambda \vec{\beta} = 0$$

$$2(X^T X + \lambda I) \vec{\beta} = 2 X^T y$$

$$\vec{\beta}_\lambda = (X^T X + \lambda I)^{-1} X^T y$$

Завдання 2. Натренуйте модель поліноміальної регресії зі степенями 1, 4, 15. Для степені 15 натренуйте модель з регуляризациєю. Спробуйте різні параметри регуляризациї. Прокоментуйте.

In [54]: `def true_fun(X):`  
 `return np.cos(1.5 * np.pi * X)`  
  
`np.random.seed(0)`  
`n_samples = 30`  
  
`X = np.sort(np.random.rand(n_samples))`  
`y = true_fun(X) + np.random.randn(n_samples) * 0.1`


In [55]: `X = np.reshape(X, (-1, 1)) # convert X to 2d array`

In [56]: `results = []`

```
In [57]: linreg = LinearRegression()
linreg.fit(X, y)
b0, b1 = linreg.intercept_, linreg.coef_[0]
y_hat = linreg.predict(X)
results.append(['LR', 1, mean_squared_error(y, y_hat)])

plt.figure(figsize=(7,4))
plt.scatter(X[:,0], y)
plt.plot(X[:,0], y_hat, color='red')
plt.title('degree = 1')
```

Out[57]: Text(0.5, 1.0, 'degree = 1')




Does not look like linear relationship at all. Simple linear regression fails.

```
In [58]: degree = 4
pf = PolynomialFeatures(degree, include_bias=False)
X_poly = pf.fit_transform(X)

linreg = LinearRegression()
linreg.fit(X_poly, y)
b0, b1 = linreg.intercept_, linreg.coef_[0]
y_hat = linreg.predict(X_poly)
results.append(['LR', degree, mean_squared_error(y, y_hat)])

plt.figure(figsize=(7,4))
plt.scatter(X[:,0], y)
plt.plot(X[:,0], y_hat, color='red')
plt.title(f'degree = {degree}')
```

Out[58]: Text(0.5, 1.0, 'degree = 4')



Looks pretty good. Line is pretty close to all points.

```

In [59]: degree = 15
pf = PolynomialFeatures(degree, include_bias=False)
X_poly = pf.fit_transform(X)

fig, axes = plt.subplots(1,2, figsize=(15, 5))
axes[0].scatter(X[:,0], y)
axes[1].scatter(X[:,0], y)

for alpha in [0, 1e-2, 1e-1, 1]:
    linreg = Lasso(alpha=alpha)
    linreg.fit(X_poly, y)
    b0, b1 = linreg.intercept_, linreg.coef_[0]
    y_hat = linreg.predict(X_poly)
    results.append((f'Lasso({alpha})', degree, mean_squared_error(y, y_hat)))

    axes[0].plot(X[:,0], y_hat, label = alpha)
    axes[0].set_title('Lasso')
    axes[0].legend(title='alpha')

for alpha in [0, 1e-2, 1e-1, 1]:
    linreg = Ridge(alpha=alpha)
    linreg.fit(X_poly, y)
    b0, b1 = linreg.intercept_, linreg.coef_[0]
    y_hat = linreg.predict(X_poly)
    results.append((f'Ridge({alpha})', degree, mean_squared_error(y, y_hat)))

    axes[1].plot(X[:,0], y_hat, label = alpha)
    axes[1].set_title('Ridge')
    axes[1].legend(title='alpha')

fig.suptitle(f'degree = {degree}')

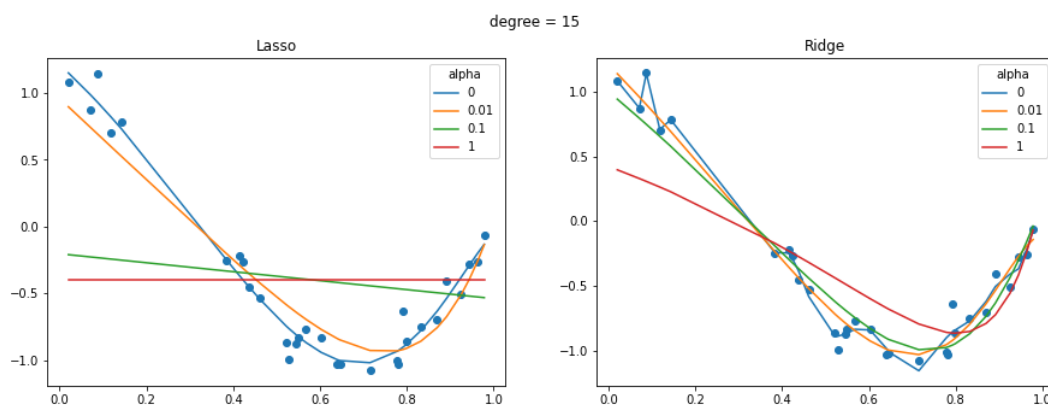
```

```

/home/daryna/anaconda3/envs/ml_ukma/lib/python3.7/site-packages/ipykernel_launcher.py:11: UserWarning: With alpha=0, this algorithm does not converge well. You are advised to use the LinearRegression estimator
# This is added back by InteractiveShellApp.init_path()
/home/daryna/anaconda3/envs/ml_ukma/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:531: UserWarning: Coordinate descent with no regularization may lead to unexpected results and is discouraged.
positive)
/home/daryna/anaconda3/envs/ml_ukma/lib/python3.7/site-packages/sklearn/linear_model/_coordinate_descent.py:531: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.17612629534727617, tolerance: 0.0012875751792313249
positive)

```

Out[59]: Text(0.5, 0.98, 'degree = 15')



Lasso(0) and Ridge(0.01) are good at predicting values. Too much regularization is bad for both models. Also Ridge(0) is overfitting - it's going to fail dramatically on new points.

```
In [60]: pd.DataFrame(results, columns=['model', 'degree', 'mse']).sort_values('mse')
```

Out[60]:

	model	degree	mse
6	Ridge(0)	15	0.004986
1	LR	4	0.011565
2	Lasso(0)	15	0.011742
7	Ridge(0.01)	15	0.012628
8	Ridge(0.1)	15	0.024254
3	Lasso(0.01)	15	0.033041
9	Ridge(1)	15	0.116239
0	LR	1	0.225892
4	Lasso(0.1)	15	0.353285
5	Lasso(1)	15	0.429192

Ridge(0) with degree=15 has smaller MSE, but it's only because of overfitting, it doesn't look like this model describes data relationships and learns the noise. Linear regression with degree 4, Lasso(0) and Ridge(0.01) have almost the same score and all of them are quite good. LR (degree=1), Lasso(0.1) and Lasso(1) with degree=15 are underfitting.

Завдання 3. Натренуйте модель логістичної регресії на вибірці про класифікацію вина. Попередньо розбийте вибірку на тренувальну та тестову частини. Яка точність отриманої моделі? Матриця невідповідностей?

```
In [61]: from sklearn.datasets import load_wine
data = load_wine()
```

```
In [62]: X, y = data['data'], data['target']
cols = data['feature_names']
```

```
In [63]: X = pd.DataFrame(X, columns = cols)
X.head()
```

Out[63]:

	alcohol	malic_acid	ash	alcalinity_of_ash	magnesium	total_phenols	flavanoids	nonflavanoid_phenol
0	14.23	1.71	2.43	15.6	127.0	2.80	3.06	0.26
1	13.20	1.78	2.14	11.2	100.0	2.65	2.76	0.28
2	13.16	2.36	2.67	18.6	101.0	2.80	3.24	0.31
3	14.37	1.95	2.50	16.8	113.0	3.85	3.49	0.20
4	13.24	2.59	2.87	21.0	118.0	2.80	2.69	0.37

```
In [64]: # train/test split
X_train, X_test, y_train, y_test = train_test_split(X, y, train_size = 0.7,
stratify=y, random_state = 42)

# standart scale data
scaler = StandardScaler()
scaler.fit(X_train)
X_train, X_test = scaler.transform(X_train), scaler.transform(X_test)
```

```
In [65]: # default params
logreg = LogisticRegression()
logreg.fit(X_train, y_train)
y_pred = logreg.predict(X_test)
print('Accuracy:', np.mean(y_test==y_pred))
confusion_matrix(y_test, y_pred)
```

Accuracy: 0.9814814814814815

```
Out[65]: array([[18,  0,  0],
               [ 1, 20,  0],
               [ 0,  0, 15]])
```

There is only one point which was wrongly predicted - point from class 1 was predicted as class 0.

```
In [66]: # optimize hyperparameters (just for fun, 98% score is already cool enough)
logreg = LogisticRegression(solver='saga',penalty='elasticnet', l1_ratio=0.0
1, C=0.1, max_iter=10000)
logreg.fit(X_train, y_train)
y_pred = logreg.predict(X_test)
print('Accuracy:', np.mean(y_test==y_pred))
```

Accuracy: 1.0

```
In [67]: confusion_matrix(y_test, y_pred)
```

```
Out[67]: array([[18,  0,  0],
               [ 0, 21,  0],
               [ 0,  0, 15]])
```

Perfect score! On this test set at least.

```
In [ ]:
```